


Categorical Imperative Programming

Type theory, refinement, and semantics for SSA

Jad-Elkhaleq Ghalayini

March 2026



This thesis is the result of my own work and includes nothing which is the outcome of work done in collaboration except as declared in the preface and specified in the text. It is not substantially the same as any work that has already been submitted, or is being concurrently submitted, for any degree, diploma or other qualification at the University of Cambridge or any other University or similar institution except as declared in the preface and specified in the text. It does not exceed the prescribed word limit for the relevant Degree Committee.

Contents

1. Introduction	4
2. S-Expressions	6
3. Evaluating Arithmetic Expressions	6
4. State	9
5. Sequencing	10
6. Branching	11
7. Loops	12
8. Unstructured Control-Flow	13
Bibliography	15
1. Sample Appendix	16

1. Introduction

The fundamental question of this thesis is: *what does a computer program mean?*

We can approach this in two ways: from the bottom-up or the top-down.

The former tends to be more natural for computer scientists.

A computer is a machine for following instructions – a program is a sequence of instructions for it to follow.

In particular, a computer has an *instruction set* – a finite set of *operations* which it can perform, which read from and write to its *state* – that is, the current state of the registers, program counters, memory, interrupts, and all the other bits which make a modern CPU tick.

Which is a *lot*.

This thesis is about the *denotational semantics* of the *static single assignment* – or *SSA* – compiler *intermediate representation* – or *IR*.

I want to start by justifying this choice of topic.

- *Why* do we need a semantics for a *compiler IR*?

Because high-level languages have too many features to effectively study *or* to lower in one pass.

Because low-level languages are designed to be *executed* efficiently – making them difficult both to reason about *or* to lower in one pass.

- *Why* should we use *SSA* as our compiler *IR*?

Practically: most modern compilers use it.

Theoretically: it's right at the intersection of functional and imperative programming – letting us use it as a *thin waist* to develop the *isomorphism* between the two.

- *Why* should we give a *denotational* semantics?

Because we want to reason about programs *algebraically* – and to do so in a *compositional* manner.

The goal of this thesis is to study the formal semantics of *SSA* – rather than do this “all at once,” we will instead take a page from the “nanopass” school of compiler development and study various intermediate languages *L*, along with the *relationships* between them.

That means that the first few chapters of this thesis are going to be a bit repetitive: for each language *L* that we introduce, after we've given some informal intuition and examples and situated it in our hierarchy of intermediate representations, we're going to need to give a formal description of the language so that we can make precise and prove our claims.

Usually, this means a variant of the following narrative:

- We give *L* a formal syntax – usually parametrized by abstract sets of operations and constants
- We'll then often formalize *L*'s “expected” behaviour by equipping it with a simple *operational semantics* – typically, we will “build up” to a full semantics for *L*, by starting with a simple rewriting-based semantics and then adding features like *state* and *nondeterminism* in stages.
- Then, we define *type theory* for *L*, consisting of:
 - A *type system* – which defines the types a term can have and the contexts in which they are interpreted. Usually, this will be parametrized by an abstract set of *base types*.

Often, we will study many different languages using the same type system.

- A *typing relation* – this typically consists of *typing rules* which tell us which terms are well-typed in a given context.
- An *equational theory* – this gives us a notion of which terms may be considered *equivalent* within a given context

At this point, we've got a formal specification for a programming language:

- How to tell whether a term is *grammatical* – like “ $x * (y + z)$ ” – or nonsense – like “ $C + +$ ”
- A collection of ways we can *interpret* a given term (natural number, string, graph, tensor, neural network...) – that is, a set of *types* A, B, C
- A description of the *contexts* in which we can interpret term – and a relation which tells us in which contexts a given term can be interpreted in a given way.

Concretely: we might be able to interpret $x + x$ as an integer in the context $x : \mathbb{Z}$, or as a string in the context $x : \text{str}$ – but $x + x$ doesn't make sense for $x : \text{POBox}$.

- An *operational semantics* which describes an *abstract machine* and tells us how to *execute* our programs on it
- An *equational theory* which tells us when two terms are “the same” in a particular context – for example, we might have $x + y \equiv y + x$ in the context $x : \mathbb{Z}, y : \mathbb{Z}$, but not in the context $x : \text{str}, y : \text{str}$.

This can let us reason about programs, but also *optimize* them – for example, we might rewrite

$$(x + 5, x + 5, x + 5) \longrightarrow \text{let } y = x + 5; 3 * y$$

Once we *have* such a specification, however, the question remains as to whether it's any *good*:

- How do we know our equational theory is *correct* – especially in the presence of subtle *potential* “gotchas” like the following optimization, which introduces UB if $n == 0$ is possible:

<pre>for(int i = 0; i < n; i++) { int t = k / n; acc += t * arr[i]; }</pre>	$\longrightarrow$	<pre>int t = k / n; for(int i = 0; i < n; i++) { acc += t * arr[i]; }</pre>
--	-------------------	--

- In fact – how do we even *define* two programs to be equivalent, in the presence of side-effects?

TODO 1: denotational semantics – do we pull the operational semantics down?

do we even want it at all –

- **it's not part of the MVP...**
- **but it really helps justify our denotational semantics as correct – otherwise completeness seems particularly abstract**

2. S-Expressions

TODO 2: This include should shift heading levels down by 1

3. Evaluating Arithmetic Expressions

We'll start by giving a formal account of the simplest practical language – that of *first-order* arithmetic expressions, like $x + y$ and $a + (5 \cdot b) + \sin(3)$.

Rather than complicate our language by introducing binary operators (like $x + y$), unary operators (like $-x$), and n -ary function applications (like $f(x, y, z)$), we can instead normalize everything to *S-expressions* – or *sexprs* – of the form $(f e_1 \dots e_n)$. We can then uniformly write:

- binary operations like $x + y$ as binary sexprs $(+ x y)$,
- unary operations $-x$ as unary sexprs $(- x)$,
- function applications $f(x, y, z)$ as n -ary sexprs $(f x y z)$.

In particular, the *operator* f always comes first, and there is no order-of-operations – all bracketing is explicit. Restricting ourselves to *first-order* sexprs means that we always require operators f to be atomic symbols drawn from a fixed set O – in particular, we don't support *partial application*, like $((\text{add } 2) 3)$, since that would require $f = (\text{add } 2)$ to be a compound expression.

Formally, we define our *syntax* as follows:

Definition 3.1 (First-order S-expressions): fixing

- a set of *variables* $x \in N$
- a set of *operations* $f \in O$
- a set of *constants* $c \in V$

we define the set of *first-order S-expressions* $e \in \text{STm}_N(O)$ to be given by the grammar

$$e ::= x \mid c \mid (f e^*)$$

We say a sexpr is *closed* if it contains no free variables – otherwise, it is *open* – where the *free variables* $\text{use}(e)$ of a sexpr e are defined as follows:

$$\text{use}(x) = \{x\} \qquad \text{use}(c) = \{\} \qquad \text{use}((f e_1 \dots e_n)) = \bigcup_i \text{use}(e_i)$$

We'll write the set of closed sexprs as $\text{STm}_\emptyset(O)$.

It's easy enough, in the language of simple arithmetic, to define what a closed sexpr “means”, since we can just *evaluate* it to find out.

For example, we might have been taught to evaluate expressions from the “inside out”, like so:

$$(+ 5 (\cdot 3 7)) = (+ 5 21) = 26$$

We can formalize this intuition in terms of a *small-step operational semantics*: given, for each operation $f \in O$, a (partial) function $\text{ev}(f) : V^* \rightarrow V$ telling us what value it returns when called with arguments $c_1, \dots, c_n$, we may define a *reduction relation* $e \rightarrow e'$ on sexprs as follows:

$$\text{step} \frac{\text{ev}(f)(c_1, \dots, c_n) = c}{(f c_1 \dots c_n) \rightarrow c} \qquad \text{arg} \frac{c_1, \dots, c_n \in V \qquad e_1 \rightarrow e'_1}{(f c_1 \dots c_n e_1 \dots e_m) \rightarrow (f c_1 \dots c_n e'_1 \dots e_m)}$$

Partiality here means that we don't need to assign a meaning to ill-defined expressions like $(/ 1 0)$ or $(+)$ (when forced, 0 is a decent choice for both of these) – we can simply leave them as undefined.

Notice that we've fixed the *evaluation order* of sexprs to be left-to-right – that is; we have to finish evaluating all terms to the *left* of a given term before we can get to reducing it.

We say a term e *evaluates* to a constant c – written $e \Downarrow c$ – if there exists a finite sequence of reductions

$$e \rightarrow e_1 \rightarrow \dots \rightarrow e_n \rightarrow c$$

i.e., if $e \rightarrow^* c$ – where R^* denotes the *reflexive, transitive closure* of a relation R .

We call a relation $e \Downarrow c$ a *big-step operational semantics* – as our original intuition for evaluation would suggest, c is a pretty good candidate for the “meaning” of a program $e \Downarrow c$.

In particular, we may define a partial function $\text{ev}(e)$ on closed terms by induction as follows:

$$\text{ev}(c) = c \qquad \text{ev}((f\ e_1 \dots e_n)) = \text{ev}(f)(\text{ev}(e_1), \dots, \text{ev}(e_n))$$

It is straightforward to show that $e \Downarrow \text{ev}(e)$ if and only if $\text{ev}(e)$ is defined.

TODO 3: Consider alternate example: $(\cdot (+ 1 1)\ x)$

- **Reduces to $(\cdot 2\ x)$ – so these should be “the same”**
- **Intuitively the same as $(+ x\ x)$**
- **Intuitively different from $(\cdot 3\ x)$**

We still, however, don't know what an *open* term means: if we try to evaluate

$$(+ (+ 3 2) (\cdot x 0)) \rightarrow (+ 5 (\cdot x 0)) \rightarrow ???$$

the first time we encounter a variable, the relation defined above becomes *stuck* – that is, we reach a term e which is not a constant, and yet does not reduce to anything – so, according to our current conventions, is ill-defined like $(/ 1 0)$.

If we fire up a Python REPL and try it, it seems to agree:

```
>>> 5 + (x * 0)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
    5 + (x * 0)
        ^
NameError: name 'x' is not defined
```

On the other hand, if x is in context – say $x = 7$ – this evaluates to 5 as we'd expect:

```
>>> x = 7
>>> 5 + (0 * x)
5
```

What this hints at is that the operational semantics of an open term e depends on the context in which it is evaluated – here, on the values of its free variables.

We can model this by defining a *contextual reduction relation* $\gamma \vdash e \rightarrow e'$ parametrized by an *evaluation context* γ – a finitely supported function $\gamma : N \rightarrow V$ assigning a values to variables – as follows:

$$\text{var} \frac{\gamma(x) = c}{\gamma \vdash x \rightarrow c} \quad \text{step} \frac{\text{ev}(f)(c_1, \dots, c_n) = c}{\gamma \vdash (f \ c_1 \dots c_n) \rightarrow c}$$

$$\text{arg} \frac{c_1, \dots, c_n \in V \quad \gamma \vdash e_1 \rightarrow e'_1}{\gamma \vdash (f \ c_1 \dots c_n \ e_1 \dots e_m) \rightarrow (f \ c_1 \dots c_n \ e'_1 \dots e_m)}$$

We can then evaluate

$$(x \mapsto 7) \vdash (+ \ (+ \ 3 \ 2) \ (\cdot \ x \ 0)) \rightarrow (+ \ 5 \ (\cdot \ x \ 0)) \rightarrow (+ \ 5 \ (\cdot \ 7 \ 0)) \rightarrow (+ \ 5 \ 0) \rightarrow 5$$

TODO 4: substitution maps open terms to closed terms

TODO 5: discuss:

- **algebra:** “equivalence under substitution”
- **programming:** “observational equivalence”

Given a set of *types* $A \in \mathcal{T}$, we’ll define a *stack-typing relation* on O to be any relation

$$f : [A_1, \dots, A_m] \rightarrow [B_1, \dots, B_n]$$

where

- $A_1, \dots, A_m \in \mathcal{T}$ are types – m is called the *input arity* or *arity* of f .
- $B_1, \dots, B_n \in \mathcal{T}$ are types – n is called the *output arity* of f .

For now, we’ll only consider operations with an output arity of 1.

We define a (*typing*) *context* $\Gamma \in \text{Ctx}_N(\mathcal{T})$ to be given by a list of *bindings*

$$x_1 : A_1, \dots, x_n : A_n$$

We can define a simple type system for S-expressions by giving a typing relation $\Gamma \vdash e : A$ that says “under the assumptions of Γ , the S-expression e has type A ”. The rules are:

TODO 6: Typing rules for constants – as distinct from nullary ops

$$\text{var-head} \frac{}{\Gamma, x : A \vdash x : A} \quad \text{var-tail} \frac{\Gamma \vdash y : B, x \neq y}{\Gamma, x : A \vdash y : B}$$

$$\text{app} \frac{f : [A_1, \dots, A_n] \rightarrow [B] \quad \forall i. \Gamma \vdash a_i : A_i}{\Gamma \vdash (f \ a_1 \dots a_n) : B}$$

TODO 7: Soundness of typing: preservation of typing

TODO 8: Completeness of typing: in the presence of simple types

TODO 9: Church vs. Curry – we will only consider well-typed terms

TODO 10: Doesn’t lose generality – untyping works in this framework due to multi-arity

TODO 11: Substitution and observational equivalence – typing rules for substitution – two open programs are equivalent iff equivalent under every substitution – this is usually where metavariables come in

TODO 12: Purity – required for substitution to preserve equivalence – typing rules for purity

TODO 13: Denotational semantics of S-expressions: partial functions

TODO 14: Denotational semantics of S-expressions: soundness and completeness

TODO 15: Equational theory of S-expressions: purity – as otherwise you can’t introduce divisions

TODO 16: Equational theory of S-expressions: soundness and completeness

4. State

TODO 17: Operational semantics of S-expressions – *state*

TODO 18: *Observational equivalence* in the presence of state

TODO 19: Denotational semantics of S-expressions – the partial state monad

TODO 20: Denotational semantics of S-expressions: *soundness* and *completeness*

TODO 21: Equational theory of S-expressions: *purity*

TODO 22: Equational theory of S-expressions: *soundness* and *completeness*

5. Sequencing

6. Branching

7. Loops

8. Unstructured Control-Flow

Bibliography

1. Sample Appendix

SAMPLE TEXT